

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.preprocessing import LabelBinarizer
```

```
In [3]: # Load the dataset
df = pd.read_csv(r"C:\Users\aksha\Downloads\Indian automobile buying behaviour study 1.0.csv")
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Age	Profession	Marrital Status	Education	No of Dependents	Personal loan	House Loan	Wife Working	Salary	Wife Salary	Total Salary	Make	Price
0	27	Salaried	Single	Post Graduate	0	Yes	No	No	800000	0	800000	i20	800000
1	35	Salaried	Married	Post Graduate	2	Yes	Yes	Yes	1400000	600000	2000000	Ciaz	1000000
2	45	Business	Married	Graduate	4	Yes	Yes	No	1800000	0	1800000	Duster	1200000
3	41	Business	Married	Post Graduate	3	No	No	Yes	1600000	600000	2200000	City	1200000
4	31	Salaried	Married	Post Graduate	2	Yes	No	Yes	1800000	800000	2600000	SUV	1600000

```
In [5]: print(df.isnull().sum())
```

```
Age                0
Profession         0
Marrital Status    0
Education          0
No of Dependents   0
Personal loan      0
House Loan         0
Wife Working       0
Salary            0
Wife Salary        0
Total Salary       0
Make              0
Price             0
dtype: int64
```

```
In [6]: # Check the data types of the columns
print(df.dtypes)
```

```
Age                int64
Profession         object
Marrital Status    object
Education          object
No of Dependents   int64
Personal loan      object
House Loan         object
Wife Working       object
Salary            int64
Wife Salary        int64
Total Salary       int64
Make              object
Price             int64
dtype: object
```

```
In [7]: # Check the distribution of the numerical columns
print(df.describe())
```

	Age	No of Dependents	Salary	Wife Salary	Total Salary \
count	99.000000	99.000000	9.900000e+01	9.900000e+01	9.900000e+01
mean	36.313131	2.181818	1.736364e+06	5.343434e+05	2.270707e+06
std	6.246054	1.335265	6.736217e+05	6.054450e+05	1.050777e+06
min	26.000000	0.000000	2.000000e+05	0.000000e+00	2.000000e+05
25%	31.000000	2.000000	1.300000e+06	0.000000e+00	1.550000e+06
50%	36.000000	2.000000	1.600000e+06	5.000000e+05	2.100000e+06
75%	41.000000	3.000000	2.200000e+06	9.000000e+05	2.700000e+06
max	51.000000	4.000000	3.800000e+06	2.100000e+06	5.200000e+06

	Price
count	9.900000e+01
mean	1.194040e+06
std	4.376955e+05
min	1.100000e+05
25%	8.000000e+05
50%	1.200000e+06
75%	1.500000e+06
max	3.000000e+06

```
In [8]: df.columns.unique()
```

```
Out[8]: Index(['Age', 'Profession', 'Marrital Status', 'Education', 'No of Dependents',
              'Personal loan', 'House Loan', 'Wife Working', 'Salary', 'Wife Salary',
              'Total Salary', 'Make', 'Price'],
              dtype='object')
```

```
In [9]: df.shape
```

```
Out[9]: (99, 13)
```

```
In [10]: df.isnull().sum().sort_values(ascending = False)
```

```
Out[10]: Age                0
Profession                0
Marrital Status          0
Education                0
No of Dependents         0
Personal loan            0
House Loan              0
Wife Working            0
Salary                  0
Wife Salary             0
Total Salary            0
Make                    0
Price                   0
dtype: int64
```

```
In [11]: # Check the distribution of the categorical columns
for col in df.select_dtypes(include=['object']):
    print(df[col].value_counts())
```

```


Profession
Salaried    64
Business    35
Name: count, dtype: int64
Marital Status
Married     84
Single      15
Name: count, dtype: int64
Education
Post Graduate    56
Graduate         43
Name: count, dtype: int64
Personal loan
No              67
Yes             32
Name: count, dtype: int64
House Loan
No              62
Yes             37
Name: count, dtype: int64
Wife Working
Yes            52
No             46
m              1
Name: count, dtype: int64
Make
SUV           19
Baleno        19
Creatia       14
i20           12
Ciaz          12
City          10
Duster         7
Verna          4
Luxuray        2
Name: count, dtype: int64

```

```

In [12]: # Histogram of the 'Age' column
plt.hist(df['Age'], bins=20, color='skyblue', edgecolor='black')
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()

```



```

In [13]: profession_column = df["Profession"]
encoder = LabelBinarizer()
one_hot_encoded_profession = encoder.fit_transform(profession_column)
df["Profession"] = one_hot_encoded_profession
print(df)

```


	Age	Profession	Marrital Status	Education	No of Dependents	\
0	27	1	Single	Post Graduate	0	
1	35	1	Married	Post Graduate	2	
2	45	0	Married	Graduate	4	
3	41	0	Married	Post Graduate	3	
4	31	1	Married	Post Graduate	2	
..	...	...	...	...	...	
94	27	0	Single	Graduate	0	
95	50	1	Married	Post Graduate	3	
96	51	0	Married	Graduate	2	
97	51	1	Married	Post Graduate	2	
98	51	1	Married	Post Graduate	2	

	Personal loan	House Loan	Wife Working	Salary	Wife Salary	Total Salary	\
0	Yes	No	No	800000	0	800000	
1	Yes	Yes	Yes	1400000	600000	2000000	
2	Yes	Yes	No	1800000	0	1800000	
3	No	No	Yes	1600000	600000	2200000	
4	Yes	No	Yes	1800000	800000	2600000	
..	...	...	...	...	...	...	
94	No	No	No	2400000	0	2400000	
95	No	No	Yes	3800000	1300000	5100000	
96	Yes	Yes	No	2200000	0	2200000	
97	No	No	Yes	2700000	1300000	4000000	
98	Yes	Yes	No	2200000	0	2200000	

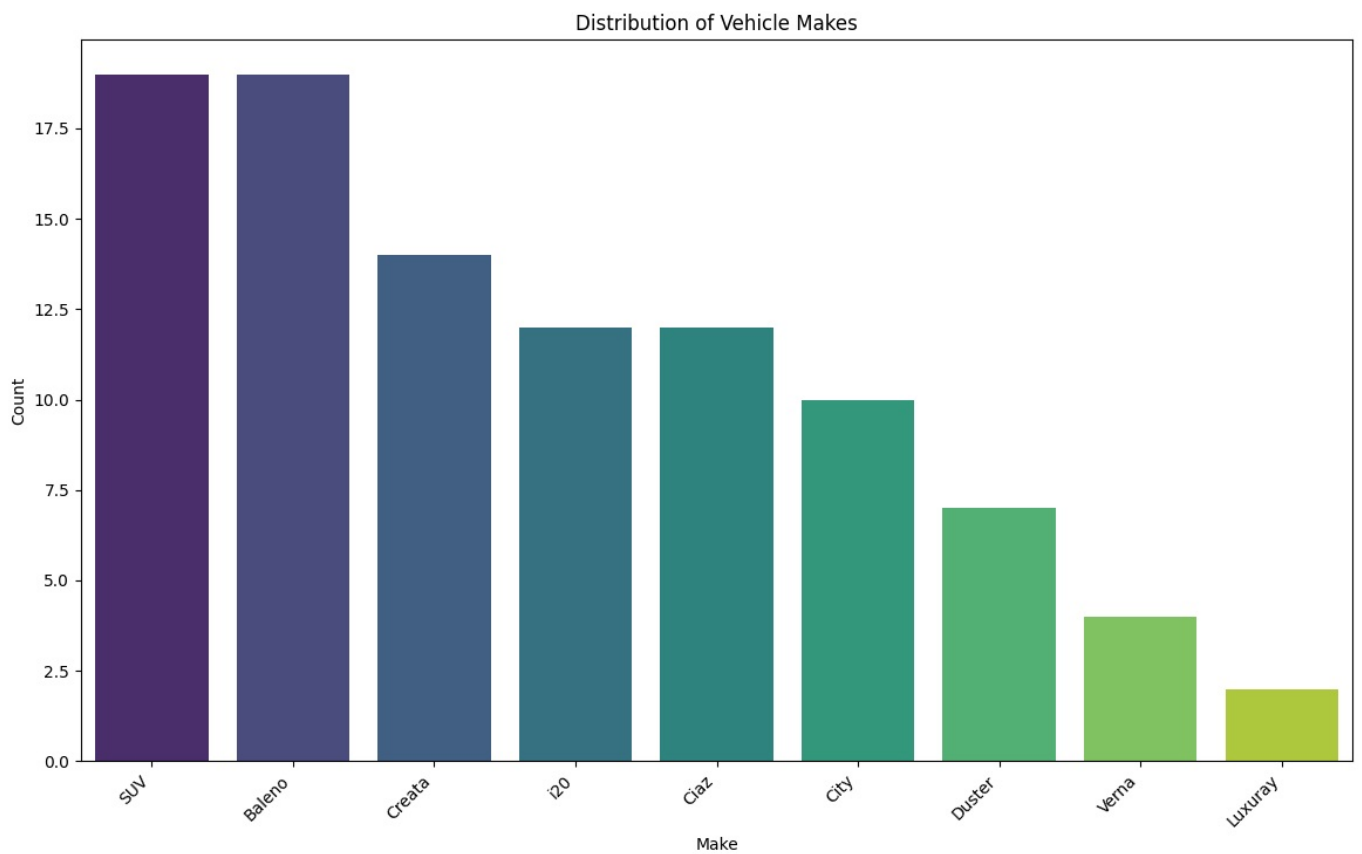
	Make	Price
0	i20	800000
1	Ciaz	1000000
2	Duster	1200000
3	City	1200000
4	SUV	1600000
..	...	...
94	SUV	1600000
95	SUV	1600000
96	Ciaz	1100000
97	Creata	1500000
98	Ciaz	1100000

[99 rows x 13 columns]

```
In [14]: # Histogram of 'Salary'
plt.hist(df['Total Salary'], bins=20, color='skyblue', edgecolor='black')
plt.title('Salary Distribution')
plt.xlabel('Salary')
plt.ylabel('Count')
plt.show()
```




```
In [15]: plt.figure(figsize=(14, 8))
sns.countplot(x='Make', data=df, order=df['Make'].value_counts().index, palette='viridis')
plt.title('Distribution of Vehicle Makes')
plt.xlabel('Make')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for better readability
plt.show()
```



```
In [16]: # Perform PCA
columns_for_pca = ['Age', 'Total Salary', 'Profession']
X = df[columns_for_pca]
scaler = StandardScaler()
X_std = scaler.fit_transform(X)
pca = PCA(n_components=3)
X_pca = pca.fit_transform(X_std)
```


```
In [17]: # Plot the PCA results
plt.scatter(X_pca[:, 0], X_pca[:, 2])
plt.xlabel('PC1')
plt.ylabel('PC3')
plt.show()
```



```
In [18]: # Perform k-means clustering
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', random_state=42)
    kmeans.fit(X_std)
    wcss.append(kmeans.inertia_)
plt.plot(range(1, 11), wcss)
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
```


```
plt.show()
```

```
C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
```



```
In [19]: # Fit the k-means model with the optimal number of clusters
kmeans = KMeans(n_clusters=3, init='k-means++', random_state=42)
y_kmeans = kmeans.fit_predict(X_std)
plt.scatter(X_pca[:, 0], X_pca[:, 2], c=y_kmeans)
plt.xlabel('PC1')
plt.ylabel('PC3')
plt.show()
```

```
C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
```




```
In [20]: def plotseg(clus,Y, labels):
plt.scatter(clus[Y==0,0], clus[Y==0,1], s=10, c='blue', label='Cluster 1')
plt.scatter(clus[Y==1,0], clus[Y==1,1], s=10, c='red', label='Cluster 2')
plt.scatter(clus[Y==2,0], clus[Y==2,1], s=10, c='black', label='Cluster 3')

plt.title('Customer_segments')
plt.xlabel(labels[0])
plt.ylabel(labels[1])
plt.show()
```

```
In [22]: clus = df.loc[:,["Age","Price"]].values
kmeans = KMeans(n_clusters=3, init='k-means++')
Y = kmeans.fit_predict(clus)
plotseg(clus, Y, ["Age","Price"])
```

C:\Users\aksha\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster\\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning  
warnings.warn(



In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js